

MAPREDUCE

Dr. Yashar Moshfeghi

Big Data 4

Tutorial, Week 2

*Part of the slides are taken from Dr. Nikos Ntarmos



Outline

- JobTracker UI
 - General information
 - Job information
 - Task information
- Hadoop programming
 - Hadoop data types
 - Basic Mapper/Reducer methods
 - Custom input/output formats
 - Word Count galore

JobTracker UI :: General information

The Jobtracker P

Details of Hadoop Installation

-version Number

-when it was compiled

-current state of the jobtracker

bigdata-06 Hadoop Map/Reduce Administration

State: RUNNING
 Started: Wed Oct 16 17:28:19 BST 2013
 Version: 2.0.0-mr1-cdh4.4.0, Unknown
 Compiled: Tue Sep 3 19:47:44 PDT 2013 by jenkins from Unknown
 Identifier: 201310161728

Cluster Summary (Heap Size is 81.06 MB/4.20 GB)

| Running Map Tasks | Running Reduce Tasks | Total Submissions | Nodes | Occupied Map Slots | Occupied Reduce Slots | Reserved Map Slots | Reserved Reduce Slots | Map Task Capacity | Reduce Task Capacity | Avg. Tasks/Node | Blacklisted Nodes | Excluded Nodes |
|-------------------|----------------------|-------------------|-------|--------------------|-----------------------|--------------------|-----------------------|-------------------|----------------------|-----------------|-------------------|----------------|
| 51 | 0 | 3 | 6 | 51 | 0 | 0 | 0 | 52 | 26 | 13.00 | 0 | 0 |

Scheduling Information

| Queue Name | State | Scheduling Information |
|------------|---------|------------------------|
| default | running | N/A |

Summary of the cluster

- measures of cluster capacity and utilization

Filter (Jobid, Priority, User, Name)

Example: 'user:smith 3200' will filter by 'smith' only in the user field and '3200' in all fields

Running Jobs

| Jobid | Priority | User | Name | Map % Complete | Map Total | Maps Completed | Reduce % Complete | Reduce Total | Reduces Completed | Job Scheduling Information | Diagnostic Info |
|-----------------------|----------|-------|-------------|-----------------------|-----------|----------------|----------------------|--------------|-------------------|----------------------------|-----------------|
| job_201310161728_0004 | NORMAL | nikos | MyWordCount | 34.33% <div></div> | 2329 | 777 | 0.00% <div></div> | 13 | 0 | NA | NA |

Failed Jobs

| Jobid | Priority | User | Name | Map % Complete | Map Total | Maps Completed | Reduce % Complete | Reduce Total | Reduces Completed | Job Scheduling Information | Diagnostic Info |
|-----------------------|----------|-------|-------------|------------------------|-----------|----------------|------------------------|--------------|-------------------|----------------------------|-----------------|
| job_201310161728_0002 | NORMAL | nikos | MyWordCount | 100.00% <div></div> | 2329 | 97 | 100.00% <div></div> | 13 | 0 | NA | NA |
| job_201310161728_0003 | NORMAL | nikos | MyWordCount | 100.00% <div></div> | 2329 | 0 | 100.00% <div></div> | 13 | 0 | NA | NA |

Retired Jobs

none

Local Logs

[Log directory](#), [Job Tracker History](#)

Hadoop, 2013.

JobTracker UI :: Job information


The Job Page

User: hdfs
 Job Name: MyWordCount
 Job Path: hdfs://bigdata-06.dcs.gla.ac.uk:8020/user/hdfs/_hadoopjob_201310161728_0004/job.xml
 Submit Host: bigdata-06.dcs.gla.ac.uk
 Submit Host Address: 132.238.255.228
 Job-ACLAR: All users are allowed
 Job Setup: Successful
 Status: Succeeded
 Started at: Wed Oct 16 17:45:54 BST 2013
 Finished at: Wed Oct 16 18:22:19 BST 2013
 Finished in: 37mins, 54secs
 Job Cleanup: Successful

| Kind | % Complete | Num Tasks | Pending | Running | Complete | Killed | Failed/Killed Task Attempts |
|--------|------------|-----------|---------|---------|----------|--------|-----------------------------|
| map | 100.00% | 2330 | 0 | 0 | 2330 | 0 | 0 / 0 |
| reduce | 100.00% | 12 | 0 | 0 | 12 | 0 | 0 / 0 |

| | Counter | | | Map | Reduce | Total |
|---|--|--|--|-------------------|----------------|-------------------|
| | | | | | | |
| File System Counters | FILE: Number of bytes read | | | 6,285,495 | 28,345 | 6,313,840 |
| | FILE: Number of bytes written | | | 375,248,934 | 2,123,757 | 380,372,691 |
| | FILE: Number of read operations | | | 0 | 0 | 0 |
| | FILE: Number of large read operations | | | 0 | 0 | 0 |
| | FILE: Number of write operations | | | 0 | 0 | 0 |
| | HDFS: Number of bytes read | | | 312,577,047,744 | 0 | 312,577,047,744 |
| | HDFS: Number of bytes written | | | 0 | 228 | 228 |
| | HDFS: Number of read operations | | | 4,722 | 6 | 4,728 |
| | HDFS: Number of large read operations | | | 0 | 0 | 0 |
| | HDFS: Number of write operations | | | 0 | 12 | 12 |
| Job Counters | Launched map tasks | | | 0 | 0 | 2,330 |
| | Launched reduce tasks | | | 0 | 0 | 12 |
| | Data-local map tasks | | | 0 | 0 | 1,291 |
| | Red-Local map tasks | | | 0 | 0 | 1,039 |
| | Total time spent by all maps in occupied slots (ms) | | | 0 | 0 | 112,029,488 |
| | Total time spent by all reduces in occupied slots (ms) | | | 0 | 0 | 5,053,251 |
| | Total time spent by all maps waiting after reserving slots (ms) | | | 0 | 0 | 0 |
| | Total time spent by all reduces waiting after reserving slots (ms) | | | 0 | 0 | 0 |
| Map-Reduce Framework | Map input records | | | 118,590,854 | 0 | 118,590,854 |
| | Map output records | | | 1,515,851,102 | 0 | 1,515,851,102 |
| | Map output bytes | | | 17,255,448,292 | 0 | 17,255,448,292 |
| | Input split bytes | | | 295,112 | 0 | 295,112 |
| | Combine input records | | | 1,515,854,983 | 37,163 | 1,515,722,126 |
| | Combine output records | | | 45,230 | 28 | 45,258 |
| | Reduce input groups | | | 0 | 12 | 12 |
| | Reduce shuffle bytes | | | 0 | 947,348 | 947,348 |
| | Reduce input records | | | 0 | 7,222 | 7,222 |
| | Reduce output records | | | 0 | 12 | 12 |
| | Spilled Records | | | 75,414 | 7,222 | 82,636 |
| | CPU time spent (ms) | | | 69,457,120 | 140,260 | 69,597,400 |
| | Physical memory (bytes) snapshot | | | 1,522,552,055,816 | 3,752,555,016 | 1,526,304,610,832 |
| | Virtual memory (bytes) snapshot | | | 4,105,586,385,512 | 23,134,134,272 | 4,128,720,519,784 |
| | Total committed heap usage (bytes) | | | 1,986,754,154,320 | 8,141,209,600 | 1,994,895,363,920 |
| UK.ac.gla.dcs.bda.WordCount\$MyMapper\$Counters | NUM_BYTES | | | 312,588,344,951 | 0 | 312,588,344,951 |
| | NUM_LINKS | | | 1,515,851,102 | 0 | 1,515,851,102 |
| | NUM_RECORDS | | | 118,590,854 | 0 | 118,590,854 |

Map Completion Graph - close



Reduce Completion Graph - close



Hadoop job_201310161728_0004 on bigdata-06

User: nikos

Job Name: MyWordCount

Job File: hdfs://bigdata-06.dcs.gla.ac.uk:8020/user/nikos/.staging/job_201310161728_0004/job.xml

Submit Host: bigdata-06.dcs.gla.ac.uk

Submit Host Address: 130.209.255.236

Job-ACLs: All users are allowed



Job Setup: Successful

Status: Running

Started at: Wed Oct 16 17:45:24 BST 2013


Running for: 8mins, 4sec

Job Cleanup: Pending

| Kind | % Complete | Num Tasks | Pending | Running | Complete | Killed | Failed/Killed Task Attempts |
|--------|---|-----------|---------|---------|----------|--------|-----------------------------|
| map | 25.87%  | 2329 | 1695 | 52 | 582 | 0 | 0 / 0 |
| reduce | 0.00%  | 13 | 13 | 0 | 0 | 0 | 0 / 0 |

| | Counter | Map | Reduce | Total |
|---|---|-------------------|--------|-------------------|
| File System Counters | FILE: Number of bytes read | 55,946 | 0 | 55,946 |
| | FILE: Number of bytes written | 113,678,798 | 0 | 113,678,798 |
| | FILE: Number of read operations | 0 | 0 | 0 |
| | FILE: Number of large read operations | 0 | 0 | 0 |
| | FILE: Number of write operations | 0 | 0 | 0 |
| | HDFS: Number of bytes read | 92,903,754,407 | 0 | 92,903,754,407 |
| | HDFS: Number of bytes written | 0 | 0 | 0 |
| | HDFS: Number of read operations | 1,460 | 0 | 1,460 |
| | HDFS: Number of large read operations | 0 | 0 | 0 |
| | HDFS: Number of write operations | 0 | 0 | 0 |
| Job Counters | Launched map tasks | 0 | 0 | 741 |
| | Data-local map tasks | 0 | 0 | 407 |
| | Rack-local map tasks | 0 | 0 | 334 |
| | Total time spent by all maps in occupied slots (ms) | 0 | 0 | 28,198,158 |
| Map-Reduce Framework | Map input records | 18,582,761 | 0 | 18,582,761 |
| | Map output records | 241,575,893 | 0 | 241,575,893 |
| | Map output bytes | 2,750,248,628 | 0 | 2,750,248,628 |
| | Input split bytes | 89,856 | 0 | 89,856 |
| | Combine input records | 240,381,830 | 0 | 240,381,830 |
| | Combine output records | 9,100 | 0 | 9,100 |
| | Spilled Records | 9,386 | 0 | 9,386 |
| | CPU time spent (ms) | 18,670,340 | 0 | 18,670,340 |
| | Physical memory (bytes) snapshot | 458,600,370,176 | 0 | 458,600,370,176 |
| | Virtual memory (bytes) snapshot | 1,236,150,968,320 | 0 | 1,236,150,968,320 |
| | Total committed heap usage (bytes) | 594,153,635,840 | 0 | 594,153,635,840 |
| uk.ac.gla.dcs.bd4.WordCount\$MyMapper\$Counters | NUM_BYTES | 92,900,417,853 | 0 | 92,900,417,853 |
| | NUM_LINES | 241,575,893 | 0 | 241,575,893 |
| | NUM_RECORDS | 18,582,761 | 0 | 18,582,761 |


Map Completion Graph - [close](#)



Reduce Completion Graph - [close](#)




Map Completion Graph - [close](#)



Reduce Completion Graph - [close](#)



Map Completion Graph - [close](#)


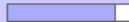





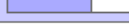








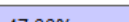
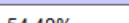
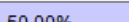
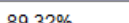
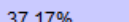
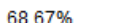


Reduce Completion Graph - [close](#)



JobTracker UI :: Task information

The Task Page

| Task | Complete | Status | Start Time | Finish Time | Errors | Counters |
|---------------------------------|---|--------|----------------------|-------------|--------|----------|
| task_201310161728_0004_m_001070 | 96.87%  | | 16-Oct-2013 18:01:38 | | | 24 |
| task_201310161728_0004_m_001074 | 94.63%  | | 16-Oct-2013 18:01:39 | | | 24 |
| task_201310161728_0004_m_001076 | 85.89%  | | 16-Oct-2013 18:01:40 | | | 24 |
| task_201310161728_0004_m_001081 | 87.50%  | | 16-Oct-2013 18:01:47 | | | 24 |
| task_201310161728_0004_m_001083 | 84.84%  | | 16-Oct-2013 18:01:48 | | | 24 |
| task_201310161728_0004_m_001085 | 84.52%  | | 16-Oct-2013 18:01:49 | | | 24 |
| task_201310161728_0004_m_001087 | 74.89%  | | 16-Oct-2013 18:01:54 | | | 24 |
| task_201310161728_0004_m_001089 | 67.75%  | | 16-Oct-2013 18:02:03 | | | 24 |
| task_201310161728_0004_m_001091 | 66.82%  | | 16-Oct-2013 18:02:04 | | | 24 |
| task_201310161728_0004_m_001094 | 79.34%  | | 16-Oct-2013 18:02:07 | | | 24 |
| task_201310161728_0004_m_001096 | 60.58%  | | 16-Oct-2013 18:02:07 | | | 24 |
| task_201310161728_0004_m_001098 | 70.95%  | | 16-Oct-2013 18:02:08 | | | 24 |
| task_201310161728_0004_m_001100 | 62.50%  | | 16-Oct-2013 18:02:10 | | | 24 |
| task_201310161728_0004_m_001102 | 64.82%  | | 16-Oct-2013 18:02:14 | | | 24 |
| task_201310161728_0004_m_001104 | 45.31%  | | 16-Oct-2013 18:02:19 | | | 24 |
| task_201310161728_0004_m_001107 | 56.77%  | | 16-Oct-2013 18:02:21 | | | 24 |
| task_201310161728_0004_m_001108 | 96.17%  | | 16-Oct-2013 18:02:24 | | | 24 |
| task_201310161728_0004_m_001109 | 47.66%  | | 16-Oct-2013 18:02:25 | | | 24 |
| task_201310161728_0004_m_001111 | 54.49%  | | 16-Oct-2013 18:02:25 | | | 24 |
| task_201310161728_0004_m_001113 | 50.00%  | | 16-Oct-2013 18:02:25 | | | 24 |
| task_201310161728_0004_m_001114 | 89.32%  | | 16-Oct-2013 18:02:26 | | | 24 |
| task_201310161728_0004_m_001115 | 37.17%  | | 16-Oct-2013 18:02:26 | | | 24 |
| task_201310161728_0004_m_001116 | 68.67% | | 16-Oct-2013 18:02:32 | | | 24 |

The Task Detail Page

Job [job_201310161728_0004](#)

All Task Attempts

| Task Attempts | Machine | Status | Progress | Start Time | Finish Time | Errors | Task Logs | Counters | Actions |
|--------------------------------------|---|---------|----------------------------------|----------------------|-------------|--------|---|----------|---------|
| attempt_201310161728_0004_m_001210_0 | /default/bigdata-06.dcs.gla.ac.uk | RUNNING | 32.47% <div><div></div></div> | 16-Oct-2013 18:04:02 | | | Last 4KB Last 8KB All | 21 | |

Input Split Locations

/default/bigdata-06.dcs.gla.ac.uk

| File System Counters | | |
|---|---------------------------------------|---------------|
| | FILE: Number of bytes read | 0 |
| | FILE: Number of bytes written | 164,566 |
| | FILE: Number of read operations | 0 |
| | FILE: Number of large read operations | 0 |
| | FILE: Number of write operations | 0 |
| | HDFS: Number of bytes read | 134,221,626 |
| | HDFS: Number of bytes written | 0 |
| | HDFS: Number of read operations | 2 |
| | HDFS: Number of large read operations | 0 |
| | HDFS: Number of write operations | 0 |
| Map-Reduce Framework | | |
| | Map input records | 31,337 |
| | Map output records | 407,381 |
| | Map output bytes | 4,637,876 |
| | Input split bytes | 128 |
| | Combine input records | 407,381 |
| | Combine output records | 13 |
| | Spilled Records | 13 |
| | CPU time spent (ms) | 33,510 |
| | Physical memory (bytes) snapshot | 660,152,320 |
| | Virtual memory (bytes) snapshot | 1,778,925,568 |
| | Total committed heap usage (bytes) | 977,600,512 |
| uk.ac.gla.dcs.bd4.WordCount\$MyMapper\$Counters | | |
| | NUM_BYTES | 134,216,895 |
| | NUM_LINES | 407,381 |
| | NUM_RECORDS | 31,337 |

Hadoop programming :: Hadoop data types

Hadoop data types

```
public interface Writable {  
    void readFields(DataInput in);  
    void write(DataOutput out);  
}
```

```
public class MyWritable implements Writable {  
    private int value;  
    private long timestamp;  
  
    public void write(DataOutput out) throws IOException {  
        out.writeInt(value);  
        out.writeLong(timestamp);  
    }  
  
    public void readFields(DataInput in) throws IOException {  
        value = in.readInt();  
        timestamp = in.readLong();  
    }  
  
    public static MyWritable read(DataInput in) throws IOException {  
        MyWritable w = new MyWritable();  
        w.readFields(in);  
        return w;  
    }  
}
```

Hadoop data types

```
public interface WritableComparable<T> extends Writable, Comparable<T> {  
    // Writable -> void readFields(DataInput in), void write(DataOutput out);  
    // Comparable<T> -> int compareTo(T o);  
}
```

```
public class MyWritableComparable implements WritableComparable {  
    private int value;  
    private long timestamp;  
  
    public void write(DataOutput out) throws IOException {  
        out.writeInt(value);  
        out.writeLong(timestamp);  
    }  
  
    public void readFields(DataInput in) throws IOException {  
        value = in.readInt();  
        timestamp = in.readLong();  
    }  
  
    public int compareTo(MyWritableComparable o) {  
        return (this.value < o.value ? -1 : (this.value == o.value ? 0 : 1));  
    }  
  
    public int hashCode() {  
        final int prime = 31;  
        return prime * (prime + value) + (int) (timestamp ^ (timestamp >>> 32));  
    }  
}
```

Hadoop data types

- ObjectWritable
- GenericWritable
- NullWritable

Generic or Abstract

- BooleanWritable
- ByteWritable
- ShortWritable, IntWritable, LongWritable
- FloatWritable, DoubleWritable
- VIntWritable, VLongWritable
- Text

Specific Java Data Types

Collection

- BytesWritable
- ArrayPrimitiveWritable, ArrayWritable, TwoDArrayWritable

- MapWritable, SortedMapWritable
- EnumSetWritable

Seldom Used

Hadoop programming :: Basic Mapper/Reducer methods

Hadoop data types

Context Object

Runs once when this mapper instance is instantiated

Called in a loop

Runs once before destructor

```
public class Mapper<KEYIN,VALUEIN,KEYOUT,VALUEOUT> {  
    static class Context { ... }  
    protected void setup(Context context) { ... }  
    protected void map(KEYIN key, VALUEIN value, Context context) { ... }  
    protected void cleanup(Context context) { ... }  
    void run(Context context) {  
        setup(context);  
  
        while (context.nextKeyValue())  
            map(context.getCurrentKey(), context.getCurrentValue(), context);  
  
        cleanup(context);  
    }  
}
```

```
public class Reducer<KEYIN,VALUEIN,KEYOUT,VALUEOUT> {  
    static class Context { ... }  
    protected void setup(Context context) { ... }  
    protected void reduce(KEYIN key, Iterable<VALUEIN> values, Context context) {  
        ...  
    }  
    protected void cleanup(Context context) { ... }  
    void run(Context context) {  
        setup(context);  
        while (context.nextKey())  
            reduce(context.getCurrentKey(), context.getValues(), context);  
        cleanup(context);  
    }  
}
```

Hadoop programming :: Custom input/output formats

Custom InputFormat

```
public abstract class InputFormat<KEY,VALUE> {  
    abstract List<InputSplit> getSplits(JobContext context);  
    abstract RecordReader<KEY,VALUE> createRecordReader(InputSplit split,  
        TaskAttemptContext context);  
}  
  
public abstract class RecordReader<KEY,VALUE> implements Closeable {  
    abstract void initialize(InputSplit split, TaskAttemptContext context);  
    abstract void close();  
    abstract boolean nextKeyValue();  
    abstract KEY getCurrentKey();  
    abstract VALUE getCurrentValue();  
    abstract float getProgress();  
}  
  
public abstract class InputSplit {  
    abstract long getLength();  
    abstract String[] getLocations();  
}
```

Hadoop programming :: Word Count galore

Word Count v0 :: Built-in mappers/reducers

Word Count v0

```
public class WordCount extends Configured implements Tool {
    public int run(String[] args) throws Exception {
        Job job = new Job();
        job.setJobName("WordCount-v0");
        job.setJarByClass(WordCount.class);

        job.setMapperClass(TokenCounterMapper.class);
        job.setReducerClass(IntSumReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.submit();
        return (job.waitForCompletion(true) ? 0 : 1);
    }

    public static void main(String[] args) throws Exception {
        System.exit(ToolRunner.run(new Configuration(), new WordCount(), args));
    }
}
```

Word Count v1 :: User-defined mappers/reducers

Word Count v1

```
public class WordCount extends Configured implements Tool {

    static class Map extends org.apache.hadoop.mapreduce.Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, Context context) throws IOException,
        InterruptedException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
        public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
        IOException, InterruptedException {
            int sum = 0;
            for (IntWritable value: values)
                sum += value.get();
            context.write(key, new IntWritable(sum));
        }
    }
}
```

Word Count v1 (cont.)

```
public int run(String[] args) throws Exception {  
    Job job = new Job(); job.setJobName("WordCount-v1");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(Map.class);  
  
    job.setCombinerClass(Reduce.class);  
    job.setReducerClass(Reduce.class);  
    job.setInputFormatClass(TextInputFormat.class);  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    job.setOutputFormatClass(TextOutputFormat.class);  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.submit();  
    return (job.waitForCompletion(true) ? 0 : 1);  
}  
  
public static void main(String[] args) throws Exception {  
    System.exit(ToolRunner.run(new Configuration(), new WordCount(), args));  
}  
}
```

The default is FileInputFormat

The default is FileOutputFormat

Word Count v2 ::
Distributed cache + configuration + counters +
status messages + progress report

Word Count v2

- Count occurrences of words in the input stream, but also:
 - Allow user to define patterns/words to be skipped
 - Count total number of words processed
 - Report progress and update status messages as we go


Word Count v2

WordCount.java

```
public class WordCount extends Configured implements Tool {
    public int run(String[] args) throws Exception {
        Job job = new Job();
        job.setJobName("WordCount-v3");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(Map.class);
        job.setCombinerClass(Reduce.class);
        job.setReducerClass(Reduce.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setOutputFormatClass(TextOutputFormat.class);

        List<String> other_args = new ArrayList<String>();
        for (int i = 0; i < args.length; ++i) {
            if ("-skip".equals(args[i])) {
                DistributedCache.addCacheFile(new Path(args[++i]).toUri(),
                                                job.getConfiguration());
                job.getConfiguration().setBoolean("wordcount.skip.patterns", true);
            } else
                other_args.add(args[i]);
        }
        FileInputFormat.setInputPaths(job, new Path(other_args.get(0)));
        FileOutputFormat.setOutputPath(job, new Path(other_args.get(1)));
    }
}
```

*Important: Configurations/
Arguments on the job client are
only propagate to mapper and
reducer if they get added here!!*



Word Count v2 (cont.)

WordCount.java (cont.)


```
        job.submit();  
        return job.waitForCompletion(true) ? 0 : 1;  
    }  
  
    public static void main(String[] args) throws Exception{  
        System.exit(ToolRunner.run(new Configuration(), new WordCount(), args));  
    }  
}
```

Same Reducer as Before


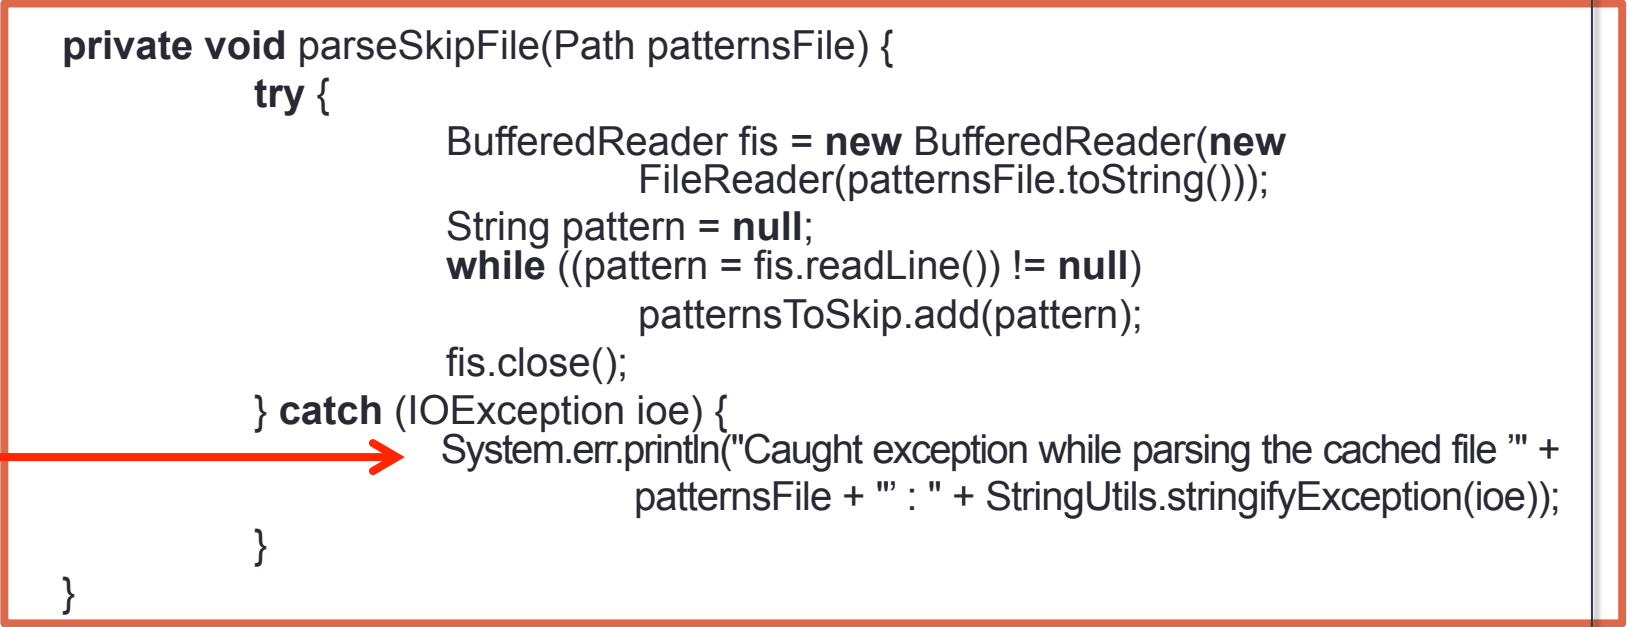
Word Count v2

Map.java

Parse the Skip File



```
public class Map extends Mapper<LongWritable, Text, Text, IntWritable>{  
    static enum Counters { INPUT_WORDS }  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
    private boolean caseSensitive = true;  
    private Set<String> patternsToSkip = new HashSet<String>();  
    private long numRecords = 0;  
    private String inputFile;
```




```
    private void parseSkipFile(Path patternsFile) {  
        try {  
            BufferedReader fis = new BufferedReader(new  
                FileReader(patternsFile.toString()));  
            String pattern = null;  
            while ((pattern = fis.readLine()) != null)  
                patternsToSkip.add(pattern);  
            fis.close();  
        } catch (IOException ioe) {  
            System.err.println("Caught exception while parsing the cached file '" +  
                patternsFile + "' : " + StringUtils.stringifyException(ioe));  
        }  
    }  
}
```

Word Count v2 (cont.)

Map.java (cont.)

```
public void setup(Context context) {
    Configuration conf = context.getConfiguration();
    inputFile = conf.get("map.input.file");

    if (conf.getBoolean("wordcount.skip.patterns", false)) {
        Path[] patternsFiles = new Path[0];
        try {
             patternsFiles = DistributedCache.getLocalCacheFiles(conf);
        } catch (IOException ioe) {
            System.err.println("Caught exception while getting cached files: " +
                               StringUtils.stringifyException(ioe));
        }
        for (Path patternsFile : patternsFiles)
            parseSkipFile(patternsFile);
    }
}

public void cleanup(Context context) {
    patternsToSkip.clear();
}
```

Word Count v2 (cont.)

Map.java (cont.)

```
public void map(LongWritable key, Text value, Context context) throws
    IOException, InterruptedException {
    String line = caseSensitive ? value.toString() :
        value.toString().toLowerCase();

    for (String pattern : patternsToSkip)
        line = line.replaceAll(pattern, "");

    StringTokenizer tokenizer = new StringTokenizer(line);
    while (tokenizer.hasMoreTokens()) {
        word.set(tokenizer.nextToken());
        context.write(word, one);
        context.getCounter(Counters.INPUT_WORDS).increment(1);
    }

    if ((++numRecords % 100) == 0)
        context.setStatus("Finished processing " + numRecords + " records " +
            "from the input file: " + inputFile);
    }
}
```

Word Count v2 (cont.)

Reduce.java

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable>
{
    public void reduce (Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable value: values)
            sum += value.get();
        context.write(key, new IntWritable(sum));
    }
}
```

Same Reducer as Before

Word Count v3 ::
Custom InputFormat + partitioner + more counters

Word Count v3

- Count occurrences of first word in each “record” in the input stream, but this time:
 - Input records are spread across multiple lines ...
 - ... with a special sequence (`\t\t\t`) separating them ...
 - Also make a custom partitioner so there are 27 reducers (for [a-z] + <everything else>)...
 - Also count total number of bytes, lines, and records processed

Word Count v3

MyPartitioner.java

```
public class MyPartitioner extends Partitioner<Text, IntWritable> {  
    int getPartition(Text key, IntWritable value, int numPartitions) {  
        int c = Character.toLowerCase(key.toString().charAt(0));  
        if (c < 'a' || c > 'z')  
            return numPartitions - 1;  
        return (int) Math.floor((float)(numPartitions - 2) * (c - 'a') / ('z' - 'a'));  
    }  
}
```

MyInputFormat.java

```
public class MyInputFormat extends FileInputFormat<LongWritable, Text> {  
    public RecordReader<LongWritable, Text> createRecordReader(InputSplit split,  
        TaskAttemptContext context) {  
        return new MyRecordReader();  
    }  
}
```


Word Count v3 (cont.)

MyRecordReader.java

```
public class MyRecordReader extends RecordReader<LongWritable, Text> {  
    private static final byte[] recordSeparator = "\t\t\t".getBytes();  
    private FSDataInputStream fsin;  
    private long start, end;  
    private boolean stillInChunk = true;  
    private DataOutputStream buffer = new DataOutputStream();  
    private LongWritable key = new LongWritable();  
    private Text value = new Text();  
  
    public void initialize(InputSplit inputSplit, TaskAttemptContext context) throws IOException {  
        FileSplit split = (FileSplit) inputSplit;  
        Configuration conf = context.getConfiguration();  
        Path path = split.getPath();  
        FileSystem fs = path.getFileSystem(conf);  
  
        fsin = fs.open(path);  
        start = split.getStart();  
        end = split.getStart() + split.getLength();  
        fsin.seek(start);  
  
        if (start != 0)  
            readRecord(false);  
    }  
}
```

Word Count v3 (cont.)

MyRecordReader.java (cont.)

```
private boolean readRecord(boolean withinBlock) throws IOException {  
    int i = 0, b;  
    while (true) {  
        if ((b = fsin.read()) == -1)  
            return false;  
        if (withinBlock)  
            buffer.write(b);  
        if (b == recordSeparator[i]) {  
            if (++i == recordSeparator.length)  
                return fsin.getPos() < end;  
        } else  
            i = 0;  
    }  
}
```

Word Count v3 (cont.)

MyRecordReader.java (cont.)

```
public boolean nextKeyValue() throws IOException {
    if (!stillInChunk)
        return false;
    boolean status = readRecord(true);
    value = new Text();
    value.set(buffer.getData(), 0, buffer.getLength());
    key.set(fsin.getPos());
    buffer.reset();
    if (!status)
        stillInChunk = false;
    return true;
}

public LongWritable getCurrentKey() { return key; }

public Text getCurrentValue() { return value; }

public float getProgress() throws IOException {
    return (float) (fsin.getPos() - start) / (end - start);
}

public void close() throws IOException { fsin.close(); }
}
```

Word Count v3 (cont.)


MyMapper.java

```
public class MyMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    static enum Counters { NUM_RECORDS, NUM_LINES, NUM_BYTES }  
    private Text _key = new Text();  
    private IntWritable _value = new IntWritable();  
  
    protected void map(LongWritable key, Text value, Context context) throws  
        IOException, InterruptedException {  
        StringTokenizer tokenizer = new StringTokenizer(value.toString(), "\n");  
        while (tokenizer.hasMoreTokens()) {  
            String line = tokenizer.nextToken();  
            int sep = line.indexOf(' ');  
            _key.set((sep == -1) ? line : line.substring(0, line.indexOf(' ')));  
            _value.set(1);  
            context.write(_key, _value);  
            context.getCounter(Counters.NUM_LINES).increment(1);  
        }  
        context.getCounter(Counters.NUM_BYTES).increment(value.getLength());  
        context.getCounter(Counters.NUM_RECORDS).increment(1);  
    }  
}
```

The diagram illustrates the execution flow of the `map` method. A red arrow originates from the `while` loop's opening brace and points to the `String line = tokenizer.nextToken();` line. Another red arrow points from the `while` loop's closing brace to the `context.getCounter(Counters.NUM_RECORDS).increment(1);` line. A third red arrow points from the `while` loop's opening brace to the `context.getCounter(Counters.NUM_RECORDS).increment(1);` line. A red rectangular box highlights the code block within the `while` loop, specifically the lines from `String line = tokenizer.nextToken();` to `context.getCounter(Counters.NUM_LINES).increment(1);`.



Word Count v3 (cont.)

MyReducer.java

```
public class MyReducer extends Reducer<Text, IntWritable, Text,
IntWritable> {
    private IntWritable _value = new IntWritable();
    protected void reduce(Text key, Iterable<IntWritable> values, Context
        context) throws IOException, InterruptedException {
        int sum = 0;
        for (Iterator<IntWritable> it = values.iterator(); it.hasNext();)
            sum += it.next().get();
        _value.set(sum);
         context.write(key, _value);
    }
}
```

Word Count v3 (cont.)

WordCount.java

```
public class WordCount extends Configured implements Tool {  
    public int run(String[] args) throws Exception {  
        Job job = new Job();  
        job.setJobName("MyWordCount(" + args[0] + ")");  
        job.setJarByClass(WordCount.class);  
         job.setInputFormatClass(MyInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
        job.setMapperClass(MyMapper.class);  
         job.setPartitionerClass(MyPartitioner.class);  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
        job.setReducerClass(MyReducer.class);  
        job.setCombinerClass(MyReducer.class);  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(job.getJobName() +  
            "_output"));  
        job.submit();  
        return job.waitForCompletion(true) ? 0 : 1;  
    }  
  
    public static void main(String[] args) throws Exception {  
        System.exit(ToolRunner.run(new WordCount(), args));  
    }  
}
```