

# HDFS

---

Dr. Yashar Moshfeghi

Big Data 4

Tutorial, Week 2

\*Part of the slides are taken from Dr. Nikos Ntarmos



# Outline

- Command-line interface
  - `hadoop fs/hdfs dfs`
- Programmatic access
  - Common parts
  - Basic operations

Command-line interface::hadoop fs/hdfs dfs

# Introduction

- There are many other interfaces to HDFS, e.g Web UI
- The command line is one of the simplest and, to many developers, the most familiar.

# HDFS CLI: Introduction

- CLI *entry points*:

\$ `hadoop fs [cmd . . .]` → **old and deprecated**

\$ `hdfs dfs [cmd ...]` → **current standard**

- These are powerful commands; use `-help` and `-usage` when in doubt!

\$ `hdfs dfs -help [cmd ...]`

*Show **long** help message for given command.*

\$ `hdfs dfs -usage [cmd ...]`

*Show just a **short message** showing usage syntax for given command.*

- Ditch the dash (-) for individual commands; e.g., for command `-mkdir`:

\$ `hdfs dfs -help -mkdir` ✗

\$ `hdfs dfs -help mkdir` ✓

# Example

\$ `hadoop fs`

Usage: `java FsShell`

```
[-ls <path>]
[-lsr <path>]
[-df [<path>]]
[-du <path>]
[-dus <path>] [-count[-q] <path>]
[-mv <src> <dst>]
[-cp <src> <dst>]
[-rm [-skipTrash] <path>]
[-rmr [-skipTrash] <path>]
[-expunge]
[-put <localsrc> ... <dst>]
[-copyFromLocal <localsrc> ... <dst>]
[-moveFromLocal <localsrc> ... <dst>]
[-get [-ignoreCrc] [-crc] <src> <localdst>]
[-getmerge <src> <localdst> [addnl]]
[-cat <src>]
[-text <src>]
[-copyToLocal [-ignoreCrc] [-crc] <src> <localdst>]
[-moveToLocal [-crc] <src> <localdst>]
[-mkdir <path>]
[-setrep [-R] [-w] <rep> <path/file>]
[-touchz <path>]
[-test [-ezd] <path>]
[-stat [format] <path>]
[-tail [-f] <file>]
[-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
[-chown [-R] [OWNER][:[GROUP]] PATH...]
[-chgrp [-R] GROUP PATH...]
[-help [cmd]]
```

# HDFS CLI: File Uploading & Downloading

- Upload file(s):

```
$ hdfs dfs -put/-copyFromLocal <file1> [<file2> ...] <dst>
```

```
$ hdfs dfs -moveFromLocal <file1> [<file2> ...] <dst>
```

- Download file(s):

```
$ hdfs dfs -get/-copyToLocal [-ignoreCrc|-crc] <file1> ... <dst>
```

```
$ hdfs dfs -moveToLocal ... (not implemented yet)
```

- Download and merge file(s):

```
$ hdfs dfs -getmerge [-nl] <src dir pattern> <dst>
```

## Example

```
$ hdfs dfs -copyFromLocal input/docs/quangle.txt hdfs://localhost/user/yashar/quangle.txt
```

// OR //

```
$ hdfs dfs -copyFromLocal input/docs/quangle.txt /user/yashar/quangle.txt
```

// OR //

```
$ hdfs dfs -copyFromLocal input/docs/quangle.txt quangle.txt
```

URI is specified  
in *core-site.xml*

Copy to home  
directory

```
$ hdfs dfs -copyToLocal quangle.txt quangle.copy.txt  
$ md5 input/docs/quangle.txt quangle.copy.txt
```

Copy back to  
local file System

```
MD5 (input/docs/quangle.txt) = a16f231da6b05e2ba7a339320e7dacd9  
MD5 (quangle.copy.txt) = a16f231da6b05e2ba7a339320e7dacd9
```



# HDFS CLI: Directory Creation & List

- Create directory:

```
$ hdfs dfs -mkdir [-p] <path>
```

- List files/directories matching pattern:

```
$ hdfs dfs -ls [-d] [-h] [-R] [<path> ...]
```

# Example

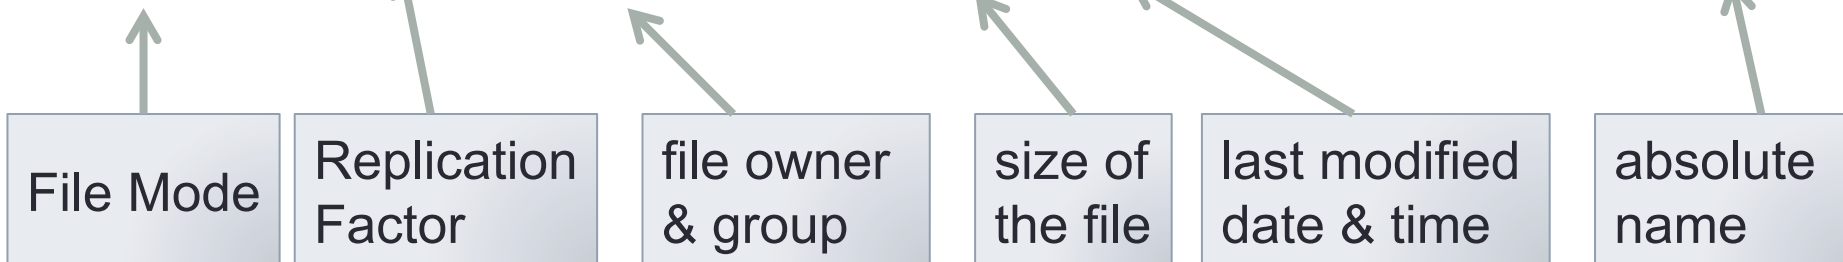
```
$ hadoop fs -mkdir books
```

```
$ hadoop fs -ls .
```

Found 2 items

drwxr-xr-x	-	yashar	supergroup	0	2009-04-02 22:41	/user/yashar/books
-rw-r--r--	1	yashar	supergroup	118	2009-04-02 22:29	/user/yashar/quangle.txt

Directories are treated as metadata and stored by the namenode, not the datanodes



traditional Unix filesystem does not have !!

# HDFS CLI: Copying, Moving

- Copy/move around a single HDFS namespace:

```
$ hdfs dfs -cp <file1> [<file2> ...] <dst>
```

```
$ hdfs dfs -mv <file1> [<file2> ...] <dst>
```

These commands allows multiple sources as well in which case the destination must be a directory.

```
$ hdfs dfs -cp /user/yashar/file1 /user/yashar/file2
```

```
$ hdfs dfs -cp /user/yashar/file1 /user/yashar/file2 /user/  
yashar/dir
```

# HDFS CLI: Deleting

- Delete HDFS file/directory:

```
$ hdfs dfs -rm [-f] [-r|-R] [-skipTrash] <pattern>
```

Only deletes non empty  
directory and files

- Delete *empty* HDFS directory:

```
$ hdfs dfs -rmdir [-ignore-fail-on-non-empty] <dir>
```

- Empty HDFS “recycle bin”:

```
$ hdfs dfs -expunge
```



*Remember !!*

Programmatic access :: Common parts

# Filesystem reference

First things first: Get a reference to the underlying filesystem.

abstract class org.apache.hadoop.fs.FileSystem

- org.apache.hadoop.fs.LocalFileSystem
- **org.apache.hadoop.fs.DistributedFileSystem**
- org.apache.hadoop.fs.HftpFileSystem
- org.apache.hadoop.fs.FTPFileSystem
- org.apache.hadoop.fs.s3.S3FileSystem
- org.apache.hadoop.fs.kfs.KosmosFileSystem

Note: In general you should strive to write your code against the FileSystem abstract class, to retain portability across filesystems !!

# Filesystem reference

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;

Configuration conf = new Configuration();
conf.set("fs.defaultFS", "hdfs://namenode.fqdn:8020"); // Not always necessary
FileSystem fs = FileSystem.get(conf);

[...]

fs.close();
```

- Configuration just a wrapper for java.util.Properties.
- Configuration's default constructor loads information from local conf files.
- Use an explicit namenode URI if executing code on a remote cluster (line 5).
- Use conf.addResource(...) to load non-standard conf (XML) files.

# Paths

- Second step: get a reference to a file/directory path
  - `org.apache.hadoop.fs.Path`
- URI-aware syntax
  - `Path f = new Path("file:///user/yashar/file.txt");`
  - `Path f = new Path("hdfs://localhost:8020/user/yashar/file.txt");`
  - `Path f = new Path("/user/yashar/file.txt");`
  - `Path f = new Path("file.txt");`



# Programmatic access :: Basic operations

# Directory/File Creation

```
import org.apache.hadoop.fs.FSDataOutputStream;  
[...]
```

```
Path dir = new Path("/path/to/dir");  
if (fs.mkdirs(dir) == false) // Create directory structure, if not there  
    // Error creating directory structure. Bail out...
```

```
Path file = new Path("/path/to/dir/file.txt");  
// ... or ...
```

```
Path file = new Path(dir, "file.txt");
```

*Absolute filename*

*Filename relative to dir !!*

```
if (fs.exists(file))  
    // File already exists. Do something...
```

```
boolean isCreated = fs.createNewFile(file);
```

```
// ... or ...
```

```
FSDataOutputStream out = fs.create(file);
```

```
// ... or ...
```

```
FSDataOutputStream out = fs.create(file, false);
```

```
// ... or ...
```

```
FSDataOutputStream out = fs.create(file, new FsPermission("u=rw,g=r,o=rwx"),  
    true, 1048576, (short)2, 134217728L, null); // ???
```

*Create empty file*

*Create/overwrite file*

*Do not overwrite file.txt*

# Input

- Read data from a file:

```
import java.io.InputStream;  
import org.apache.hadoop.fs.FSDataInputStream;  
[...]
```

```
InputStream in = fs.open(file);  
// ... or ...  
FSDataInputStream in = fs.open(file);  
String str = in.readUTF();  
int i = in.readInt();  
double d = in.readDouble();  
  
in.close();
```

- Also check out `org.apache.commons.io.IOUtils`.

# Output

- Write data to a file:

```
import org.apache.hadoop.fs.FSDataOutputStream;  
[...]
```

```
FSDataOutputStream out = fs.create(file);  
out.writeUTF("1024");  
out.writeInt(1024);  
out.writeDouble(1.0);
```

```
out.close();
```

# Example

```
public class FileSystemDoubleCat {  
  
    public static void main(String[] args) throws Exception {  
        String uri = args[0];  
        Configuration conf = new Configuration();  
        FileSystem fs = FileSystem.get(URI.create(uri), conf);  
        FSDataInputStream in = null;  
        try {  
            in = fs.open(new Path(uri));  
            IOUtils.copyBytes(in, System.out, 4096, false);  
        }  
        finally { IOUtils.closeStream(in);  
        }  
    }  
}
```

Here's the result of running it on a small file:

```
% hadoop FileSystemDoubleCat hdfs://localhost/user/yashar/quangle.txt
```

On the top of the Crumpetty Tree  
The Quangle Wangle sat,  
But his face you could not see,

# Seeking

- Moving around within a file only supported for reading

```
FSDatInputStream in = fs.open(file);  
in.seek(1024); // Go to position 1024 (in bytes)  
String someString = in.readUTF(); // Read a string  
in.skip(1024); // Skip next 1024 bytes  
int position = in.getPos(); // Get current position  
in.seek(0); // Go back to start
```

- HDFS built for streaming access ⇒ Seeking not a good idea **(expensive)**!
- Seeking not available for writes ⇒ Only append is supported...

```
FSDatOutputStream out = fs.append(file);  
out.writeUTF("some text");
```

... but not by all FileSystem subclasses!

# Example

```
public class FileSystemDoubleCat {  
  
    public static void main(String[] args) throws Exception {  
        String uri = args[0];  
        Configuration conf = new Configuration();  
        FileSystem fs = FileSystem.get(URI.create(uri), conf);  
        FSDataInputStream in = null;  
        try {  
            in = fs.open(new Path(uri));  
            IOUtils.copyBytes(in, System.out, 4096, false);  
            in.seek(0); // go back to the start of the file  
            IOUtils.copyBytes(in, System.out, 4096, false);  
        }  
        finally { IOUtils.closeStream(in);  
        }  
    }  
}
```

Here's the result of running it on a small file:

```
% hadoop FileSystemDoubleCat hdfs://localhost/user/yashar/quangle.txt
```

On the top of the Crumpey Tree  
The Quangle Wangle sat,  
But his face you could not see,

---

On account of his Beaver Hat.  
On the top of the Crumpey Tree The Quangle Wangle sat,  
But his face you could not see, On account of his Beaver Hat.

# Directory/file renaming and deletion

```
Path dir = new Path("/path/to/dir");  
Path file = new Path(dir, "file.txt");
```

*// Move file.txt to .../other/dir and rename it*

```
Path dst = new Path("/path/to/other/dir/otherfile.txt");  
fs.rename(file, dst);
```

*// Delete the new file*

```
fs.delete(dst);
```

*// Delete the original directory recursively*

```
fs.delete(dir, true);
```



# Further Reading

- Apache Hadoop Documentation
  - User Guide:  
<http://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-hdfs/HdfsUserGuide.html>
  - Architecture:  
<http://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>
  - Command-Line Interface:  
<http://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-hdfs/HDFSCommands.html>
  - Java API:  
<http://hadoop.apache.org/docs/stable/api/index.html>  
(look under org.apache.hadoop.fs)
- Cloudera Library
  - <http://www.cloudera.com/content/cloudera/en/documentation.html#ClouderaDocumentation>
- Google . . .